

Lecture 2: Optimization methods

The top of the slide features a blue header with a background of binary code (0s and 1s) and a magnifying glass icon. The word "Overview" is written in white, centered in the header.


Overview

- Introduction
- Mathematical formulation
- One-dimensional minimization
- Newton's Method
- Conjugate gradient method

Introduction

Optimization is used everywhere in nature:

- Physics
- Chemistry
- Economics
- Engineering
- Etc...



and of course **STATISTICS!**

Introduction

- Example 1: Industry

How to produce a cylindrical beer can 0.5L so it requires minimum material?

Continuous optimization



Introduction

- Example 2: Economics

Factories F1, F2

Retail outlets R1, R2, R3

Cost of shipping a product c_{ij}

Production a_i each week

Requirement b_j each week

Network flow optimization



F1




F2



Introduction

Example 3: Statistics

Maximize Likelihood $L(X, \theta)$



- Almost all model fitting requires optimization!

Maximum likelihood

Consider a sample (X_1, \dots, X_n) which is drawn from a probability distribution $P(X|\Theta)$ where Θ are parameters.

If the X s are independent with probability density function $P(X_i|\Theta)$ then the joint probability of the whole set is

$$P(X_1, \dots, X_n / \Theta) = \prod_{i=1}^n P(X_i / \Theta)$$

Find the parameters that maximize this function

Mathematical formulation

We need to minimize or maximize

- **Objective function** $f(x)$ (I - cost, II - profit, III-likelihood)

dependent on

- **Parameters** or **Unknowns** x (I-height & diameter, II-supply, III – parameters)

Mathematical formulation

- Sometimes we have constraints $c_i(x)$ satisfying equations or inequalities.
Formulation:

$$\min_{x \in R^n} f(x) \text{ subject to } \begin{aligned} c_i(x) &= 0, i \in E \\ c_i(x) &\geq 0, i \in I \end{aligned}$$

What if:

- Max instead of min
- Constraints are not like these

Mathematical formulation

- Example 1: Constraints – volume=0.5L
- Example 2- cont.

$$\begin{aligned} \min \quad & \sum_{ij} c_{ij} x_{ij} \\ & \sum_{j=1}^3 x_{ij} \leq a_i, i = 1, 2 \\ \text{s.t.} \quad & \sum_{i=1}^2 x_{ij} \geq b_j, j = 1, 2, 3 \\ & x_{ij} \geq 0 \end{aligned}$$

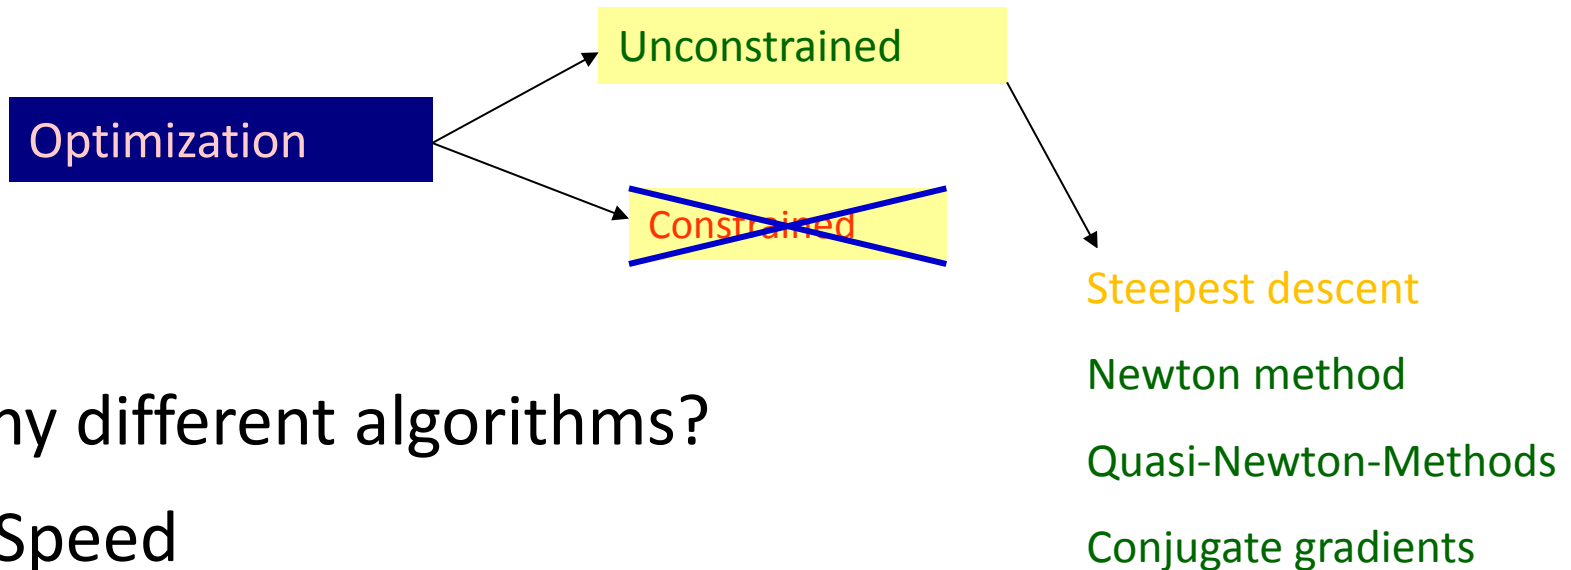
- Example 3 – no constraints **UNCONSTRAINED MINIMIZATION**



Exercise

- Split into groups of three-four and
 1. Find an application when optimization is needed (your personal experience, research, university courses)
 2. State your problem
 - Objective function
 - Parameters
 - Constraints if any
 3. You have **max 10 minutes**

Where we are



Why different algorithms?

- Speed
- Memory
- Historically


One-dimensional minimization

- One-dimensional minimization=one parameter
- Algorithm Golden Section: finds local minimum on interval $[A,B]$
- It narrows down the search interval, constant reduction factor $1-\alpha=(\sqrt{5}-1)/2\approx 0.62$

One-dimensional minimization

Golden section

1. Choose interval $[x_1, x_3]$
2. Choose $a = \alpha(x_3 - x_1)$
3. $x_2 = x_1 + a$, $x_4 = x_3 - a$
4. If $f_4 > f_2$ select RED
5. If $f_4 < f_2$ select BLUE
6. Continue with new interval until it is small



Note: f should be unimodal

R: One-dimensional minimization

- Brent's method – improved golden search

`optimize(f, interval,...)`

Multidimensional optimization

The problem:

$$\min_{\mathbf{x} \in R^n} f(\mathbf{x})$$

Gradient

$$\nabla f(\mathbf{x}) = \begin{pmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} \\ \dots \\ \frac{\partial f(\mathbf{x})}{\partial x_n} \end{pmatrix}$$

Hessian

$$\nabla^2 f(\mathbf{x}) = \begin{pmatrix} \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_1} & \dots & \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_n} \\ \dots & \dots & \dots \\ \frac{\partial^2 f(\mathbf{x})}{\partial x_n \partial x_1} & \dots & \frac{\partial^2 f(\mathbf{x})}{\partial x_n \partial x_n} \end{pmatrix}$$

General methodology:

1. Given starting point x_0 , $x=x_0$
2. Choose direction p and step α
3. Move to $x:=x+ \alpha p$
4. Repeat from 2 until convergence

Multidimensional optimization

- How to choose direction leading to function decrease ?

Taylor theorem $f(x + \alpha p) = f(x) + \alpha p^T \nabla f(x_k) + o(\alpha^2)$

The minimum is $p = \frac{-\nabla f(x)}{\|\nabla f(x)\|}$

Should be minimized

Any direction having $\angle(d, -\nabla f(x)) < \frac{\pi}{2}$ is descent direction

Multidimensional optimization

- How to choose step size α ?
 - Find global minimum along direction p (expensive)
 - Find a sufficient decrease

BACKTRACKING

Choose $\alpha_0 > 0$, ρ in $(0,1)$, c in $(0,1)$, $\alpha := \alpha_0$

REPEAT until $f(x_k + \alpha p_k) \leq f(x_k) + c \alpha \nabla f_k^T(p_k)$

$\alpha := \rho \alpha$

END

Newton's method

- In statistics called Newton-Raphson method

General idea:

Quadratic model

$$f(\mathbf{p}) = \frac{1}{2} \mathbf{p}^T A \mathbf{p} + b^T \mathbf{p} + c$$

Minimum

$$\mathbf{p}^* = A^{-1}b$$

- When general function, Taylor expansion

$$f(x + \alpha p) \approx f(x) + \alpha \nabla f(x)^T p + \frac{\alpha^2}{2} p^T \nabla^2 f(x) p$$

Proceed to next point


$x := x + \alpha p$

$$p = -(\nabla^2 f(x))^{-1} \nabla f(x)$$

Newton's method

Illustration:

- Steepest descent
- Newton's direction



Newton's method

Comments

- Under mild conditions: Converges quickly, especially near optimum
- For p to be a descent direction Hessian should be **positive definite** (see why)—strong requirement!
- Can be very expensive to compute reverse of Hessian on each iteration!
- Need to store $n \times n$ matrix (Hessian) – memory requirements

Quasi-Newton methods

Idea:

In Newton's method instead computing inverse of Hessian on each step

- Compute approximate Hessian B_k and reverse H_k
- **BFGS**: Using knowledge about H_k , function and gradient in x_k and x_{k+1} , compute H_{k+1}

$$p_k = -H_k \nabla f(x_k)$$

BFGS

How to compute H_{k+1} ?

- Quadratic model

$$m_{k+1}(p) = f(x) + \nabla f(x)^T p + \frac{1}{2} p^T B_{k+1} p$$

should have the same function values and gradients as $f(x)$ in points x_k and x_{k+1}

-> Secant condition

$$H_{k+1} (\nabla f_{k+1} - \nabla f_k) = x_{k+1} - x_k$$

y_k s_k

BFGS

How to compute H_{k+1} ?

- Distance between H_k and H_{k+1} should be minimal

$$\min_H \|H - H_k\|$$

$$s.t. H = H^T, \text{ secant condition}$$

- Updating formula

$$H_{k+1} = H_k - \frac{H_k y_k y_k^T H_k}{y_k^T H_k y_k} + \frac{s_k s_k^T}{y_k^T s_k}$$

BFGS

Comments

- Typically takes more iterations than Newton's method
- Each iteration takes less time (no matrix inversion!)
- Quasi-Newton Methods are particularly good for large-scale problems.
- How to choose initial Hessian?

Conjugate Gradient method

Quadratic function

$$f(x) = \frac{1}{2} x^T A x - b^T x$$

Gradient

$$\nabla f(x) = Ax - b \stackrel{\text{def}}{=} r(x)$$

A- symmetric, positive definite

Def. Directions p and q are conjugate with respect to A if

$$p^T A q = 0$$

Conjugate Gradient method

Conjugate gradient method:

Choose

$$p_{k+1} = -r_{k+1} + \beta_{k+1}p_k$$

$$p_0 = -r_0$$

p_i should satisfy conjugacy condition, therefore

- $\beta_k = \frac{r_k^T A p_{k-1}}{p_{k-1}^T A p_{k-1}}$
- Converges in $\dim(A)$ steps

Nonlinear CG method

Idea: Consider general $f(x)$ and substitute r_k with ∇f_k

Given $x_0, f_0, \nabla f_0$

$p_0 := -\nabla f_0$

while $\nabla f_k \neq 0$

 compute $\alpha_k, x_{k+1} = x_k + \alpha_k p_k$

$\beta_{k+1} = (\nabla f_{k+1}^T \nabla f_{k+1}) / (\nabla f_k^T \nabla f_k)$

$p_{k+1} = -\nabla f_{k+1} + \beta_{k+1} p_k$

$k = k+1$

end

Nonlinear CG method

- Converges to local minimum
- Much faster than steepest descent in general
- Slower than Newton and Quasi-Newton but much less memory

R: Multidimensional optimization

- Quasi-Newton and CG incorporated in one procedure

`optim(par, fn, gr=NULL, method, ...)`

- Look also

`nls(...)`